HANDS-ON MACHINE LEARNING (HOME)

Lecture/Lab – 2

An Introduction to Machine Learning Using LLMs as an Example

Goal: Run an LLM (Mistral 7B or Llama 2) on a local machine and chat with it

Ghulam Rasool, PhD

Department of Machine Learning, Moffitt Cancer Center

OVERVIEW

- **Review of LLMs**
- **Running LLMs on Local Machines**
  - LM Studio
  - Ollama
- **Graphical User Interface**
  - Open WebUI
  - Docker
- **Llama.cpp and gguf**
- **Prompting and Retrieval Augmented Generation (RAG)**
  - LangChain
  - LlamaIndex

LARGE LANGUAGE MODELS (LLMs)

Brief History

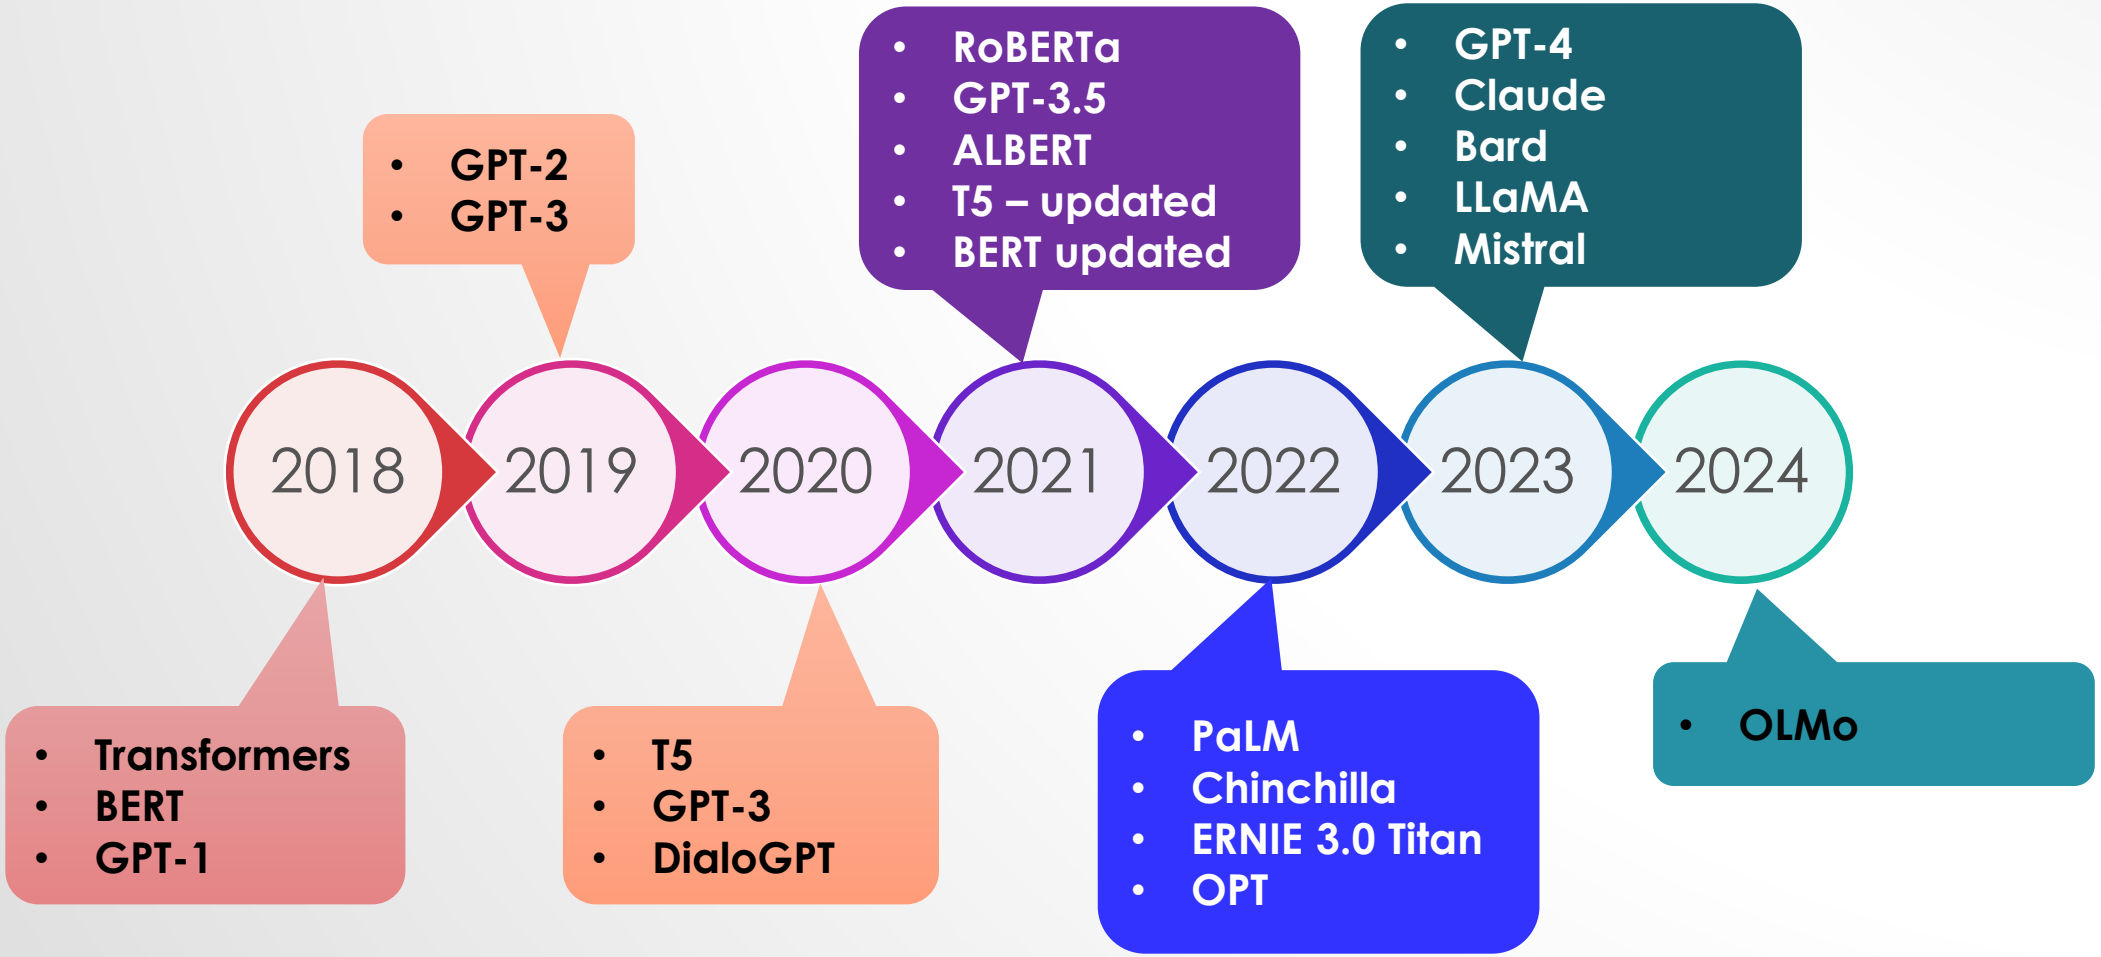

LARGE LANGUAGE MODELs (LLMs)

Large

- **Higher capacity to learn**

Billion of variables to store the learned knowledge

- **Larger datasets**

Trillions of bytes of data to learn from

LARGE LANGUAGE MODELs (LLMs)

How to train your LLM

Data

LLM

GPUs

Millions $

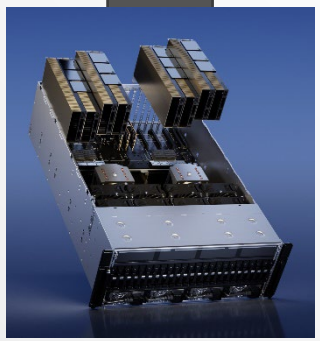

LARGE LANGUAGE MODELs (LLMs)

What does training mean?

- Predict the next word

She is moving to Tampa for her job.

Gradient Descent Algorithm

She is moving to Tampa for her job.

She is moving to Tampa for her family.

- Correct – Pat
- Incorrect – Punish

She is moving to Tampa for her car.

She is moving to Tampa for her alligator.

LARGE LANGUAGE MODELs (LLMs)

As a result of the training, we have

She is moving to Tampa for her _____.

What is the most probable next word?

Job 49%

Family 49%

Car 0.5%

Alligator 1.5%

LARGE LANGUAGE MODELs (LLMs)

As a result of the training, we have

She is moving to Tampa for her _____.

What is the most probable next word?

Job 49%

Family 49%

Car 0.5%

Alligator 1.5%

Prompt Engineering

Sabrina is a family person. She is moving to Tampa for her _____.

What is the most probable next word?

Job 9%

Family 89%

Car 0.5%

Alligator 1.5%

Hallucination

Temperature – 0 to 1

LARGE LANGUAGE MODELs (LLMs)

How to train your LLM

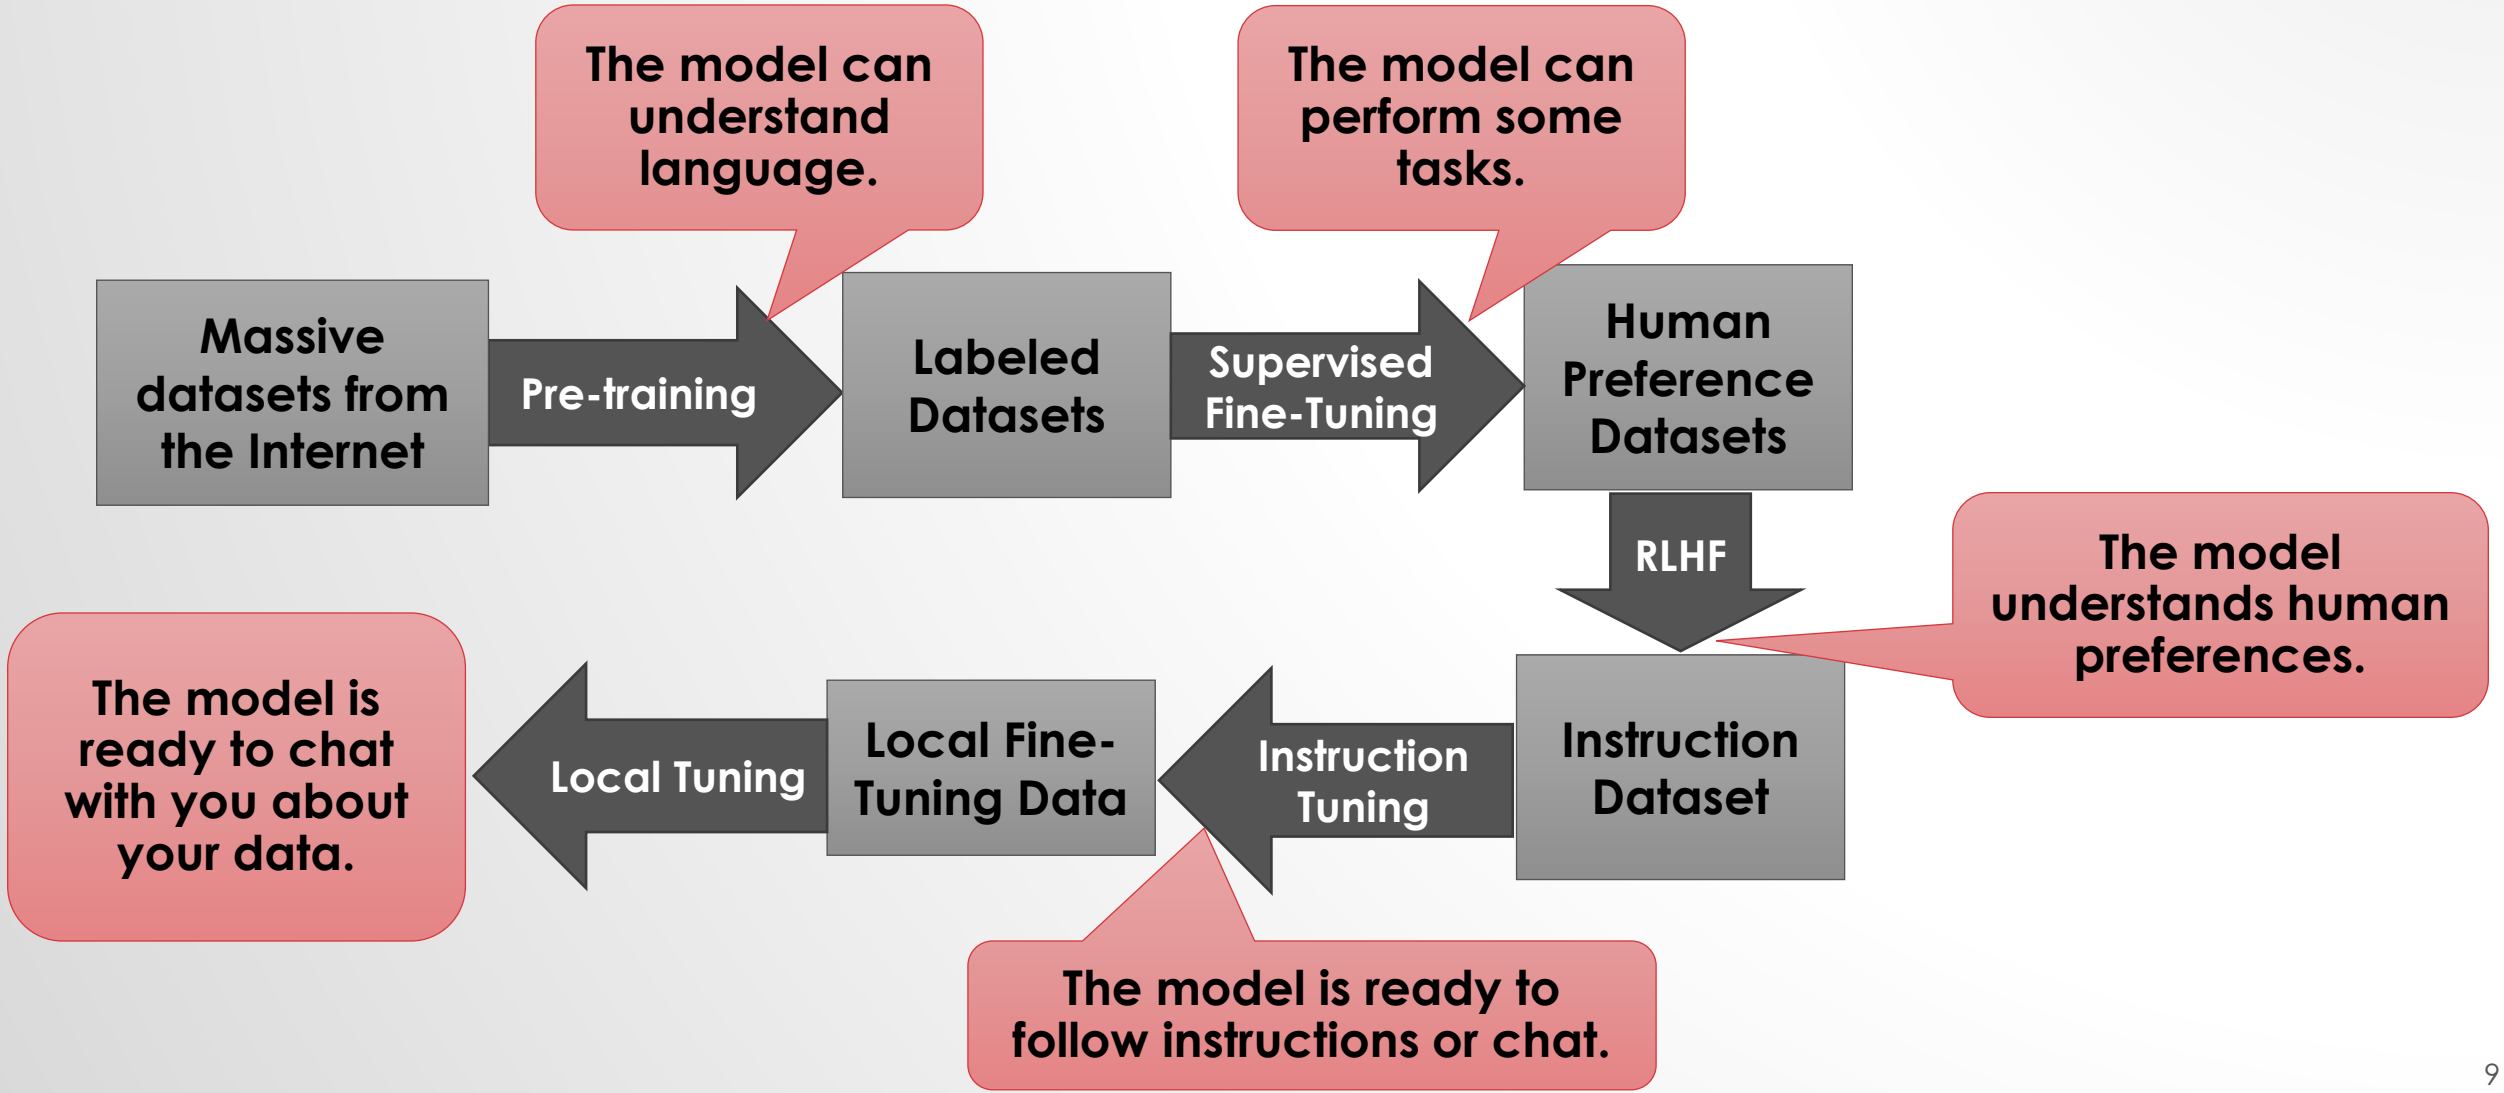

LARGE LANGUAGE MODELs (LLMs)

How to train your LLM

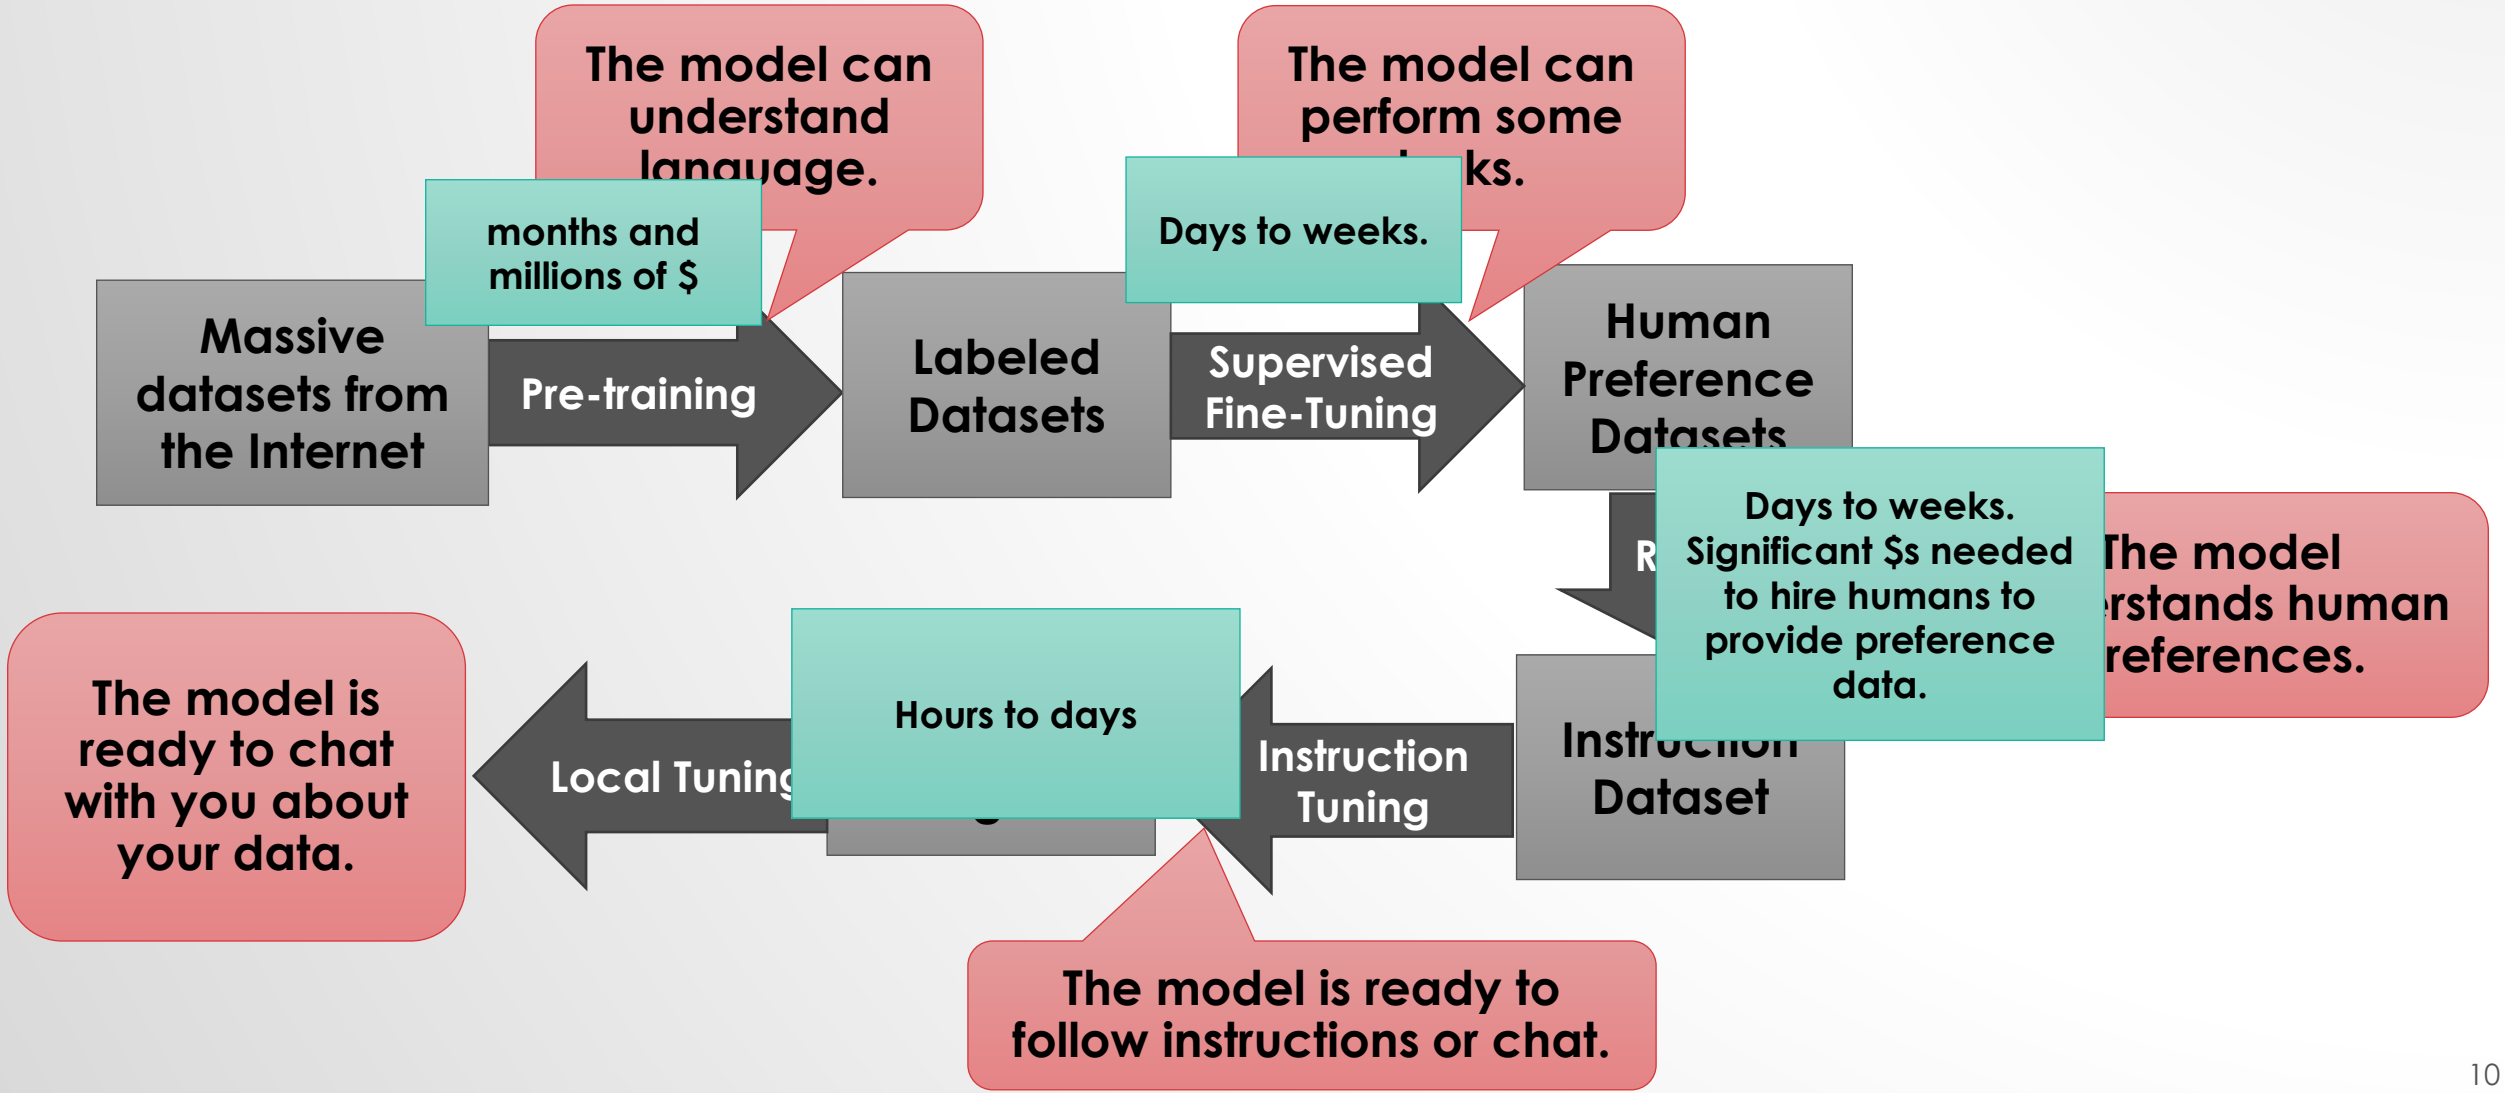

Chemotherapy is the standard of care in cancer treatment.

Ch	20394
##em	5521
##otherapy	20939
is	1110
the	1103
standard	2530
of	1104
care	1920
in	1107
cancer	4182
treatment	3252
.	119

```
Embedding: tensor([-4.0514e-01,  3.2718e-01,  2.3717e-01,  3.2862e-01,  3.6550e-01,
         2.1226e-02,  6.9171e-01,  7.1098e-02, -9.7321e-03, -7.7804e-01,
        -8.0872e-01, -1.8547e-01, -5.6361e-02,  4.9584e-01,  1.3878e-01,
         6.7449e-01,  3.7221e-01, -7.5023e-01, -9.9605e-02, -3.0082e-01,
        -6.1156e-02, -6.8750e-02,  2.5605e-01, -2.1948e-01,  1.1133e-01,
        -3.5326e-01,  3.8680e-01,  7.2568e-01, -5.1890e-01, -2.8369e-01,
        -2.8277e-02,  3.2367e-01,  3.0425e-01,  9.0662e-01, -9.2147e-01,
         6.3722e-01,  8.4916e-01, -1.2166e-01,  6.7056e-01, -3.6339e-01,
        -1.4453e-01,  5.5456e-01, -2.1817e-01, -4.5966e-01,  4.9136e-02,
        -3.3405e-01, -1.9264e-01, -2.6836e-01,  9.8193e-02, -9.1319e-01,
         3.9529e-01,  6.3396e-01,  1.0694e+00, -4.0752e-01, -1.1956e-01,
         8.3672e-01, -3.2265e-01,  4.9057e-02, -1.8049e-01,  1.2337e-01,
         1.1135e+00, -2.1958e-01,  5.2144e-01, -3.5725e-02,  7.7396e-01,
        -2.3286e-01, -1.0921e+00, -1.0853e-01, -1.2074e+00, -4.5416e-02,
        -1.0770e-01, -8.7412e-02,  4.6003e-01,  1.7978e-01, -4.7101e-01,
        -2.9541e-01, -2.0189e-01,  2.7894e-01,  3.8826e-01, -2.8794e-01,
        -3.1304e-01,  3.1930e-01,  4.5050e-02,  1.0765e+00,  7.8386e-01,
        -7.8647e-01, -7.8930e-02,  5.2840e-02, -3.3437e-01,  7.0197e-01,
        -5.4875e-01,  4.5861e-02, -4.2728e-01, -4.2825e-01,  4.4960e-03,
         3.8803e-01,  1.3139e-01,  1.1247e-01, -3.1398e-01, -3.5722e-01,
         4.6070e-01, -1.7379e-01,  2.8147e-01, -7.9178e-01, -2.5676e-01,
         5.5337e-02,  2.9408e-01, -5.4813e-01,  2.5966e-02, -2.5847e-01,
        -6.7750e-01,  3.4987e-01, -5.6569e-01, -3.4727e-02,  6.8431e-02,
        -1.2239e-01,  4.4732e-01, -3.6277e-01, -1.3723e-01, -2.1545e-01,
...
         1.1014e-01, -7.5589e-01, -4.8585e-01, -6.5725e-01, -7.1004e-01,
         1.9709e-01,  3.6595e-01,  4.9644e-01])

768
```

LARGE LANGUAGE MODELs (LLMs)

Model Training

All or part of the model parameters are updated

- Transfer Learning
- Unsupervised Learning
- Supervised Learning
- Curriculum Learning

Adapter Training

Original model parameters are frozen

- Add additional small models, called adapters
- Train “adapters” for individual tasks tasks.
- Low-Rank Adaptation
- Parameter-Efficient Fine-Tuning

Prompt Engineering

Model parameters are frozen

- No adapters are added
- Add more context to the prompt
- Zero-shot
- Few-shot
- Retrieval Augmented Generation (RAG)

LARGE LANGUAGE MODELs (LLMs)

Traditional NLP/ML

- Needs labelled data
  - Cost of data collection/labeling
  - Legal/Privacy concerns around using data
- 1 model per task results in
  - Increased model development/tuning cost
  - Increased operational costs
  - Increased money spent on sourcing data
- Relatively Limited generalization
- Computationally cheaper (~300 Million parameters)

Modern NLP/ML

- Need a small amount of labeled data
- A single generic model can do more than one task
- More generalized: Besides language, it learns higher-level concepts, styles, etc.
- Computationally Expensive (~500 Billion parameters)

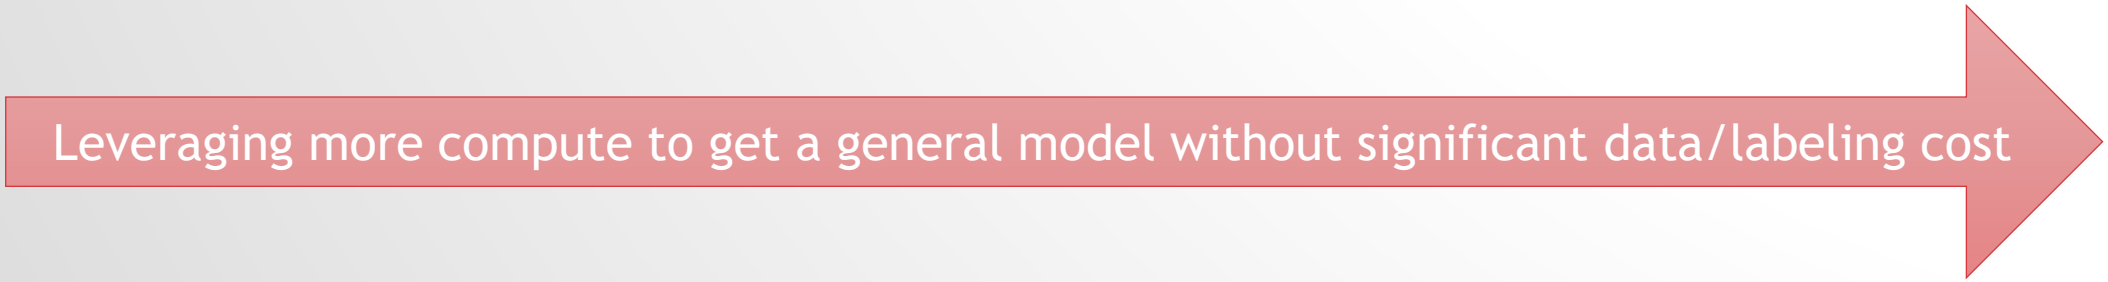

Running LLMs on Local Machines

- LM Studio
- Ollama
- Open WebUI

LM STUDIO

- **Run LLMs on local computers entirely offline**
- **Use models through the in-app Chat UI or an OpenAI-compatible local server**
- **Download any compatible model files from Hugging Face repositories**

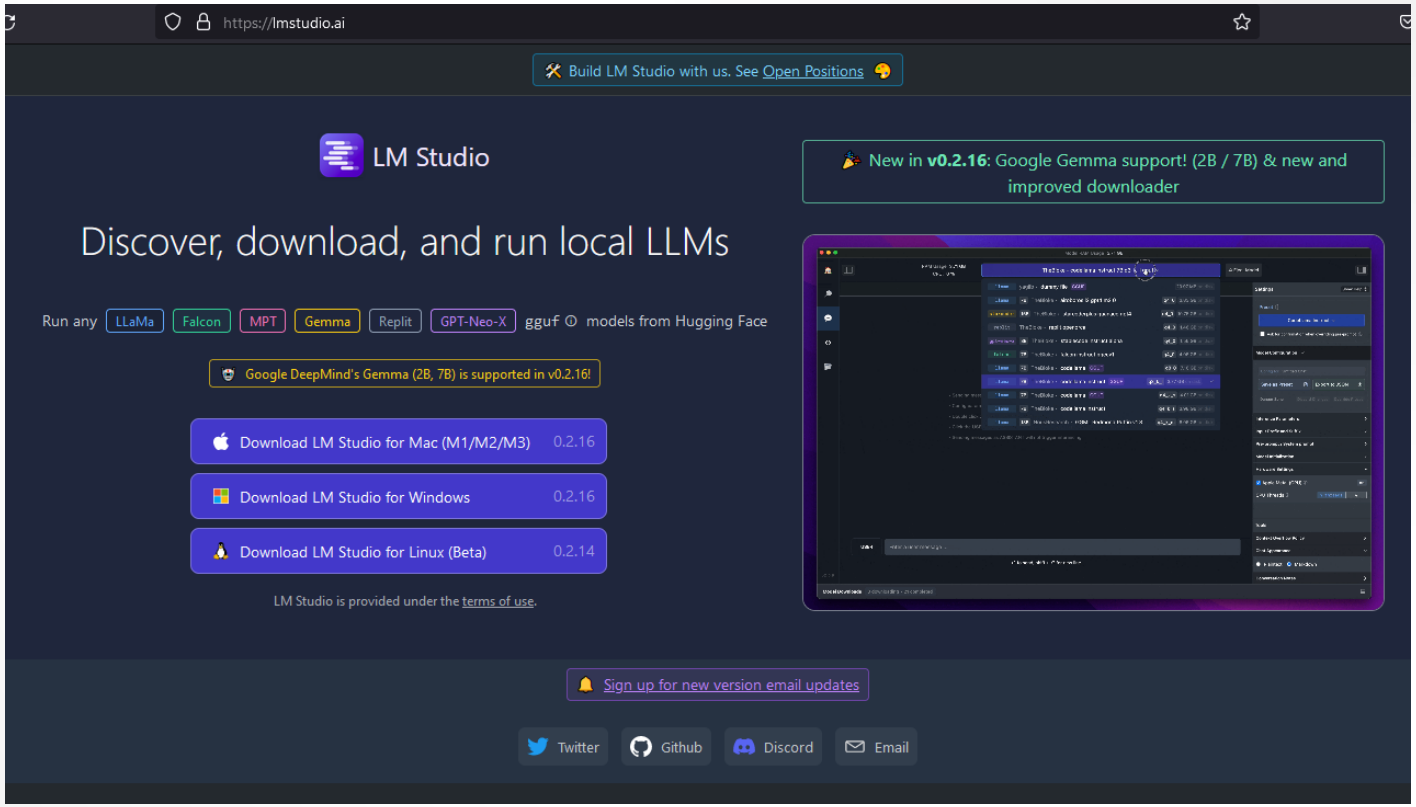

LM STUDIO

- **Quantization**
  - **Quantization refers to a set of techniques that enable running models on resource-constrained platforms**
  - **Q2_K, Q4_K_M, Q5_0, Q8_0**
- **Chat UI**
- **GPU Offload**
- **Context Length**

https://towardsdatascience.com/introduction-to-weight-quantization-2494701b9c0c

LM STUDIO - LOCAL INFERENCE SERVER

- **Serve the model on the local machine**
  - http://localhost:1234/
  - http://127.0.0.1:1234
- **How do I use it?**

```python
# Example: reuse your existing OpenAI setup
from openai import OpenAI

# Point to the local server
client = OpenAI(base_url="http://localhost:1234/v1", api_key="not-needed")

completion = client.chat.completions.create(
  model="local-model", # this field is currently unused
  messages=[
    {"role": "system", "content": "Always answer in rhymes."},
    {"role": "user", "content": "Introduce yourself."}
  ],
  temperature=0.7,
)

print(completion.choices[0].message.content)
```

✓ 2.4s

```
In verse and rhyme, I'm here to entertain,
A friendly AI, with a poetic brain.
Delighted to meet you, may our chat remain,
A joyful exchange of words, so calm and sane.
```

https://towardsdatascience.com/introduction-to-weight-quantization-2494701b9c0c

LM STUDIO - LOCAL INFERENCE SERVER

- **Serve the model on the local machine**
  - http://localhost:1234/
  - http://127.0.0.1:1234
- **How do I use it?**
  - **Chatbot**
  - **VLM Chatbot**
    - **Llava Base Model**
    - **Vision Adapter**

https://towardsdatascience.com/introduction-to-weight-quantization-2494701b9c0c

OLLAMA

https://ollama.com/

- Running LLMs, locally

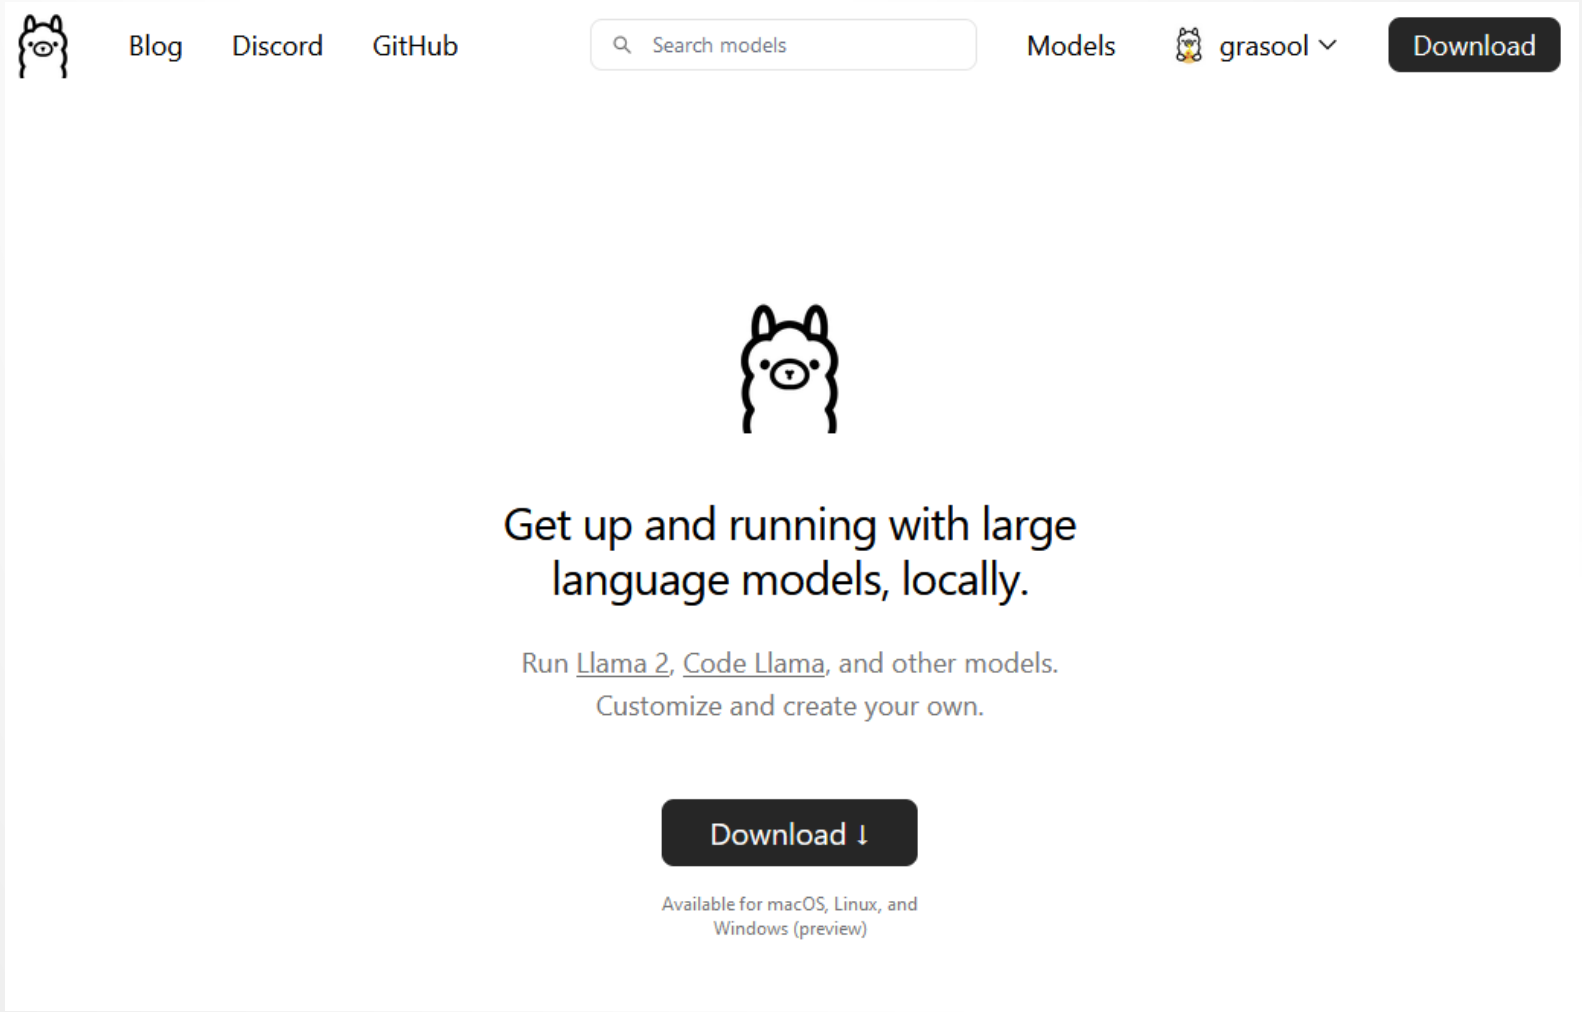

OLLAMA

https://ollama.com/

- Running LLMs, locally
- http://localhost:11434

OPEN WEBUI

https://openwebui.com/

https://github.com/open-webui/open-webui

- http://localhost:3000

- Download and Install Docker

https://docs.docker.com/desktop/install/windows-install/